


Classic ML vs. Fine-tuned DistilBERT

A comparison of test Macro F1 scores



What the chart shows

The fine-tuned DistilBERT model scores 0.821 on test Macro F1, compared with 0.664 for the Classic ML model. Its score is higher by 0.157 F1 points.

Macro F1 averages each class's F1 score equally. It reflects performance across classes rather than allowing a more common class to dominate the average.

The chart reports one test score for each model. It does not show class-level results, sample size, or uncertainty, so it cannot establish how consistent the difference would be across other datasets or runs.